

## Homework 7. Replication in Cassandra

### 1) Configure 3 node cluster

Commands were taken from [this tutorial](#). As a docker image, I used [bitnami/cassandra](#).

I used a docker-compose located in the [repo](#) to create a network, run docker containers, and expose ports.

```
(base) denys_herasymuk@EPUALVIW07D6:~/UCU/UCU_DE_Program_2022_2023/Distributed_Databases/UCU_DE_Distributed_Databases/HW7_Cassandra_Replication$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
4852d9b16d6e	bitnami/cassandra:4.1.0	"/opt/bitnami/script..."	2 minutes ago	Up 2 minutes	7000/tcp, 9042/tcp
36d1e82ddab4	bitnami/cassandra:4.1.0	"/opt/bitnami/script..."	2 minutes ago	Up 2 minutes	7000/tcp, 9042/tcp
551ee1a4c26a	bitnami/cassandra:4.1.0	"/opt/bitnami/script..."	2 minutes ago	Up 2 minutes	7000/tcp, 9042/tcp

### 2) Validate configuration

```
(base) denys_herasymuk@EPUALVIW07D6:~/UCU/UCU_DE_Program_2022_2023/Distributed_Databases/UCU_DE_Distributed_Databases/HW7_Cassandra_Replication$ docker exec -ti cassandra-node1 nodetool status
```

```
Datacenter: datacenter1
=====
Status=Up/Down
-- State=Normal/Leaving/Joining/Moving
-- Address      Load       Tokens     Owns (effective)  Host ID                               Rack
UN 172.27.0.4    74.73 KiB   256        66.4%             aea92e49-9700-4a98-a8c2-298c6b6bd1c4 rack1
UN 172.27.0.3    74.7 KiB    256        65.8%             ec8a8316-bb59-48e7-a5fd-a6e72131a614 rack1
UN 172.27.0.2    74.7 KiB    256        67.8%             0597ebf7-9cb3-4e51-9538-8ee4372e085d rack1
```

### 3) Create 3 keyspaces with replication factor 1, 2, 3 using *cqlsh*

Connect to the node 1 and run *cqlsh*.

```
$ docker exec -it cassandra-node1 bash
I have no name!@5749f4941419:/$ cqlsh
```

Create three keyspaces with a replication factor 1, 2, 3, respectively.

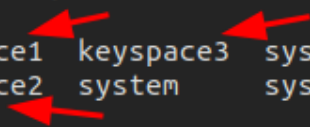
```
cassandra@cqlsh> CREATE KEYSPACE keyspace1 WITH REPLICATION = {'class' :
'NetworkTopologyStrategy', 'replication_factor' : 1};

cassandra@cqlsh> CREATE KEYSPACE keyspace2 WITH REPLICATION = {'class' :
'NetworkTopologyStrategy', 'replication_factor' : 2};

cassandra@cqlsh> CREATE KEYSPACE keyspace3 WITH REPLICATION = {'class' :
'NetworkTopologyStrategy', 'replication_factor' : 3};
```

```
cassandra@cqlsh> DESCRIBE keyspaces
```

```
keyspace1  keyspace3  system_auth      system_schema    system_views  
keyspace2  system      system_distributed  system_traces    system_virtual_schema
```



#### 4) Create a table in each of the keyspaces

```
CREATE TABLE keyspace1.user_location (  
    id int primary key,  
    name text,  
    country text  
);
```

```
CREATE TABLE keyspace2.user_location (  
    id int primary key,  
    name text,  
    country text  
);
```

```
CREATE TABLE keyspace3.user_location (  
    id int primary key,  
    name text,  
    country text  
);
```

#### 5) Write and read to / and from different nodes

Write to and read from keyspace 1 on the node 1.

```
INSERT INTO keyspace1.user_location (id, name, country) VALUES (1, 'User1', 'Ukraine');
```

```
SELECT * FROM keyspace1.user_location;
```

```
cassandra@cqlsh> SELECT * FROM keyspace1.user_location;
```

```
id | country | name  
----+-----+-----  
1  | Ukraine | User1  
(1 rows)
```

Write to and read from keyspace 1 on the node 2.

```
$ docker exec -it cassandra-node2 bash
```

```
cqlsh> INSERT INTO keyspace1.user_location (id, name, country) VALUES (2, 'User2', 'Poland');
```

```
cqlsh> SELECT * FROM keyspace1.user_location;
```

```
cassandra@cqlsh> SELECT * FROM keyspace1.user_location;
```

id	country	name
1	Ukraine	User1
2	Poland	User2

(2 rows)

6) Insert the data into the created tables and look at their distribution across the cluster nodes

Execute the following commands for each of the keyspaces.

```
INSERT INTO keyspace1.user_location (id, name, country) VALUES (3, 'User3', 'Ukraine');
INSERT INTO keyspace1.user_location (id, name, country) VALUES (4, 'User4', 'Poland');
...
INSERT INTO keyspace1.user_location (id, name, country) VALUES (9, 'User9', 'Ukraine');
INSERT INTO keyspace1.user_location (id, name, country) VALUES (10, 'User10', 'Poland');
```

Look at the data distribution.

```
I have no name!@91e74d56e8ce:/$ nodetool status keyspace1;
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
-- Address      Load          Tokens   Owns (effective)  Host ID                               Rack
UN  172.29.0.3    75.18 KiB     256      33.7%             8fcbbcda-3894-49e4-b58f-e7dd7c944c6b rack1
UN  172.29.0.2    75.21 KiB     256      32.3%             6ccdf7f1-c27c-4520-9d94-e6fa3fa358f3 rack1
UN  172.29.0.4    75.19 KiB     256      34.0%             2fb0d6f2-4860-4dd0-872c-4db88f0f4d76 rack1
```

```
I have no name!@91e74d56e8ce:/$ nodetool status keyspace2;
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
-- Address      Load          Tokens   Owns (effective)  Host ID                               Rack
UN  172.29.0.3    75.18 KiB     256      66.1%             8fcbbcda-3894-49e4-b58f-e7dd7c944c6b rack1
UN  172.29.0.2    75.21 KiB     256      67.7%             6ccdf7f1-c27c-4520-9d94-e6fa3fa358f3 rack1
UN  172.29.0.4    75.19 KiB     256      66.2%             2fb0d6f2-4860-4dd0-872c-4db88f0f4d76 rack1
```

```
I have no name!@91e74d56e8ce:/$ nodetool status keyspace3;
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
-- Address      Load          Tokens   Owns (effective)  Host ID                               Rack
UN 172.29.0.3    75.18 KiB     256      100.0%           8fcbbcda-3894-49e4-b58f-e7dd7c944c6b rack1
UN 172.29.0.2    75.21 KiB     256      100.0%           6ccdf7f1-c27c-4520-9d94-e6fa3fa358f3 rack1
UN 172.29.0.4    75.19 KiB     256      100.0%           2fb0d6f2-4860-4dd0-872c-4db88f0f4d76 rack1
```

7) For any record from each of the keyspaces, display the nodes on which the data is stored

```
I have no name!@91e74d56e8ce:/$ nodetool getendpoints keyspace1 user_location 1
172.29.0.3
I have no name!@91e74d56e8ce:/$ nodetool getendpoints keyspace2 user_location 1
172.29.0.3
172.29.0.2
I have no name!@91e74d56e8ce:/$ nodetool getendpoints keyspace3 user_location 1
172.29.0.3
172.29.0.2
172.29.0.4
```

8) Disable one of the nodes. For each of the keyspaces, determine with which levels of *consistency* we can read and write, and which of them provide *strong consistency*

Disable node 3.

```
(base) denys_herasymuk@EPUALVIN0706:~$ docker stop cassandra-node3
cassandra-node3
(base) denys_herasymuk@EPUALVIN0706:~$ docker ps
CONTAINER ID   IMAGE                  COMMAND                  CREATED        STATUS        PORTS                               NAMES
91e74d56e8ce   bitnami/cassandra:4.1.0 "/opt/bitnami/script..." About an hour ago Up About an hour 7000/tcp, 0.0.0.0:9043->9042/tcp, :::9043->9042/tcp cassandra-node2
5749f4941419   bitnami/cassandra:4.1.0 "/opt/bitnami/script..." About an hour ago Up About an hour 7000/tcp, 0.0.0.0:9042->9042/tcp, :::9042->9042/tcp cassandra-node1
(base) denys_herasymuk@EPUALVIN0706:~$
```

Consistency levels in Cassandra

<https://docs.datastax.com/en/cassandra-oss/3.0/cassandra/dml/dmlConfigConsistency.html>

Execute the following commands to understand the allowed read consistency levels for each keyspace.

```
CONSISTENCY <level>;

SELECT * FROM keyspace1.user_location LIMIT 3;
SELECT * FROM keyspace2.user_location LIMIT 3;
SELECT * FROM keyspace3.user_location LIMIT 3;
```

Example of the output.

```
cqlsh> CONSISTENCY ONE;
Consistency level set to ONE.
cqlsh> SELECT * FROM keyspace1.user_location LIMIT 3;
NoHostAvailable: ('Unable to complete the operation against any hosts', {<Host: 127.0.0.1:9042 data center1>: Unavailable('Error from server: code=1000 [Unavailable exception] message="Cannot achieve consistency level ONE" info={\'consistency\': \'ONE\', \'required_replicas\': 1, \'alive_replicas\': 0}\')})})
cqlsh> SELECT * FROM keyspace2.user_location LIMIT 3;
```

id	country	name
1	Ukraine	User1
2	Poland	User2
3	Ukraine	User3

(3 rows)

```
cqlsh> SELECT * FROM keyspace3.user_location LIMIT 3;
```

id	country	name
1	Ukraine	User1
2	Poland	User2
3	Ukraine	User3

## Read consistency summary

- **Keyspace 1:** None of the consistency levels works.
- **Keyspace 2:** Works only with ONE.
- **Keyspace 3:** Works with ONE, TWO, SERIAL, QUORUM (but not THREE or ALL).

Execute the following commands to understand the allowed write consistency levels for each keyspace.

```
CONSISTENCY <level>;

INSERT INTO keyspace1.user_location (id, name, country) VALUES (11, 'User11', 'Ukraine');
INSERT INTO keyspace2.user_location (id, name, country) VALUES (11, 'User11', 'Ukraine');
INSERT INTO keyspace3.user_location (id, name, country) VALUES (11, 'User11', 'Ukraine');
```

## Write consistency summary

- **Keyspace 1:** Works with ANY, ONE.
- **Keyspace 2:** Works with ANY, ONE.
- **Keyspace 3:** Works with ANY, ONE, TWO, QUORUM (but not THREE or ALL).

## Strong consistency summary

Strong consistency can be achieved if  $W + R > RF$ , where  $R$  – read CL replica count,  $W$  – write CL replica count,  $RF$  – replication factor ([src](#)).

- **Keyspace 1:** Since there is only one replica, read and write operations are strongly consistent..
- **Keyspace 2:** No strongly consistent combinations.
- **Keyspace 3:** The following combinations are strongly consistent ( $R + W$ ): TWO + TWO, TWO + QUORUM, QUORUM + QUORUM, SERIAL + TWO, SERIAL + QUORUM.

## 9-11) Create network partition and value conflict

Write different values on each of the nodes.

```
# ===== Write one record only on node 1 =====
docker stop cassandra-node2
docker stop cassandra-node3

cassandra-node1: INSERT INTO keyspace3.user_location (id, name, country) VALUES (20,
'User20', 'Ukraine');

# ===== Write one record only on node 3 =====
docker start cassandra-node3
docker stop cassandra-node2

cassandra-node3: INSERT INTO keyspace3.user_location (id, name, country) VALUES (20,
'User2000', 'Ukraine3');

# ===== Write one record only on node 2 =====
docker start cassandra-node2
docker stop cassandra-node1

cassandra-node2: INSERT INTO keyspace3.user_location (id, name, country) VALUES (20,
'User200', 'Ukraine2');

docker start cassandra-node1
docker start cassandra-node2
```

First, reading with **consistency level ONE** returned id=20, name=User20, country=Ukraine. However, in several seconds it returned id=20, name=User200, country=Ukraine2.

Read with **consistency level QUORUM/ALL** returns id=20, name=User200, country=Ukraine2.

This means that Read Repair occurred, and the nodes synchronized and used the last written value for the id.

```
cqlsh> CONSISTENCY QUORUM;
Consistency level set to QUORUM.
cqlsh> SELECT * FROM keyspace3.user_location WHERE id = 20;

id | country | name
---+-----+-----
20 | Ukraine2 | User200

(1 rows)
```

```
cqlsh> CONSISTENCY ALL;
Consistency level set to ALL.
cqlsh> SELECT * FROM keyspace3.user_location WHERE id = 20;

id | country | name
---+-----+-----
20 | Ukraine2 | User200

(1 rows)
```

```
cqlsh> CONSISTENCY ONE;
Consistency level set to ONE.
cqlsh> SELECT * FROM keyspace3.user_location WHERE id = 20;

id | country | name
---+-----+-----
20 | Ukraine2 | User200

(1 rows)
```

## 12) Check the behavior of lightweight transactions for previous points in a non-partitioned and partitioned cluster

### 1. Non-partitioned cluster case

```
cqlsh> UPDATE keyspace3.user_location SET name = 'New User' WHERE id = 20 IF name = 'User201';

[applied] | name
-----+-----
      False | User200

cqlsh> UPDATE keyspace3.user_location SET name = 'New User' WHERE id = 20 IF name = 'User200';

[applied]
-----
      True

cqlsh> SELECT * FROM keyspace3.user_location WHERE id = 20;

id | country | name
---+-----+-----
 20 | Ukraine2 | New User

(1 rows)
```

### 2. Partitioned cluster case

Disable node 2 and 3.

```
(base) denys_herasymuk@EPUALVIW07D6:~$ docker ps
CONTAINER ID   IMAGE                      COMMAND                  CREATED        STATUS
PORTS          NAMES
56b75149a974   bitnami/cassandra:4.1.0   "/opt/bitnami/script..." About an hour ago Up 16 minutes
7000/tcp, 0.0.0.0:9042->9042/tcp, :::9042->9042/tcp   cassandra-node1
(base) denys_herasymuk@EPUALVIW07D6:~$
```

```
cqlsh> UPDATE keyspace3.user_location SET name = 'New New User' WHERE id = 20 IF name = 'New User';
NoHostAvailable: ('Unable to complete the operation against any hosts', {<Host: 127.0.0.1:9042 data center1>: Unavailable('Error from server: code=1000 [Unavailable exception] message="Cannot achieve consistency level SERIAL" info={\'consistency\': \'SERIAL\', \'required_replicas\': 2, \'alive_replicas\': 1}\')})
cqlsh>
```

Since lightweight transactions have SERIAL consistency level (similar to QUORUM), the update query above requires at least 2 available nodes.